

Capstone 1 - Personalized Portfolio Advisor with What-If Simulation

- Create a goal-based advisory agent where users enter financial goals, risk appetite, and current holdings, and receive a suggested allocation. The agent should stress-test the portfolio against scenarios (rate hike, market downturn, inflation), visualize projected outcomes in charts, and suggest rebalancing steps.
- Suitability disclaimers must be included. Bonus: multi-agent roles (risk profiler, allocator, simulator).
- Objective: Deliver personalized, scenario-tested investment guidance with clear risk communication.
- Learning takeaways: Goal-based financial planning logic, Monte Carlo / scenario simulation, data visualization, risk profiling, and embedding responsible-advice guardrails and disclaimers.

Agent	Assigned To	Description
Data Agent	Tushar Bansal	Ingests the portfolio CSV, fetches live prices via yfinance and news headlines per ticker; outputs a clean, structured dataset for everything downstream
Analysis Agent	Rajeev Sahu	Computes portfolio metrics — returns, Sharpe ratio, sector allocation, concentration — and produces the analytics report the advisory chain reasons from
Risk Agent	Nagjyothi Tammaana	Scores the investor's risk profile (conservative / moderate / aggressive) and evaluates whether current holdings match that tolerance
Advisory Agent	Syed Hussein	Generates investment and rebalancing suggestions grounded in the RAG knowledge layer, tailored to the risk profile
QA Agent	Shivani Khobrekar	The compliance and suitability gate — checks every recommendation before it reaches the user, and can send it back for revision

Capstone 2 - Conversational Financial Statement Analyzer

- Create a GenAI tool where users upload company financials (P&L, balance sheet, cash flow) and ask questions conversationally. It should compute key ratios (liquidity, leverage, profitability), detect trends across periods, highlight red flags, and produce a plain-English narrative summary with citations back to the source figures.
- Bonus: compare two companies side by side.
- Objective: Turn raw financial statements into grounded, queryable insight without manual spreadsheet work.
- Learning takeaways: Retrieval-augmented generation (RAG) over documents, tool-use for reliable numeric computation, controlling hallucination on financial figures, grounding/citation techniques, and financial- ratio interpretation.

Agent	Assigned To	Description
Trend & Red Flag Agent	Shubham Shukla	Computes period-over-period deltas across all ratios and evaluates the six red-flag patterns (accrual-quality gap, DSO climbing, margin compression, leverage stress, inventory build-up, liquidity deterioration)
Retrieval / Grounding Agent	Suman Kumar	Indexes normalized statements and narrative notes in ChromaDB; on each question, assembles the grounded context block of computed facts plus relevant text
Narrative & Q&A Agent	Abhishek Tripathi	The conversational surface — answers from grounded context only, with a mandatory citation on every figure, and declines when a number isn't in the statements
Comparison Agent	Ram M Singh	Runs two companies side by side, surfacing fiscal-year-end or currency mismatches rather than comparing regardless

Capstone 3 - Intelligent Invoice & Expense Reconciliation Agent

- Build an agentic system that ingests invoices, receipts, and bank statements (PDF/image/CSV), performs a three-way match against purchase orders, and flags mismatches, duplicates, and anomalies. The agent should explain each exception in plain language and route uncertain cases to a human for approval.
- Bonus: auto-draft a follow-up email to the vendor for disputed items.
- Objective: Automate transaction reconciliation and surface only the exceptions that need human attention.
- Learning takeaways: Document parsing and OCR, structured data extraction from messy inputs, entity/transaction matching logic, anomaly detection, and designing human-in-the-loop approval flows.

Agent	Assigned To	Description
Ingestion Agent	Roshan Javvaji	Accepts a mixed folder of PDFs, images and CSVs; routes by type (text-layer PDF, OCR for scans, pandas for CSV) and emits tagged raw text and tables
Extraction Agent	Jayashri Sheeli	Turns messy raw text into strict Pydantic records, handling missing fields, varied date formats and mangled OCR tables — attaches a per-field confidence score
Matching Agent	Sambhaji S Pandhare	Performs the three-way match (PO ↔ invoice ↔ receipt) with configurable tolerances and fuzzy vendor matching. Deterministic Python, no LLM — this has to be auditable
Anomaly Agent	Shivam Thaman	Independent of matching: detects duplicate invoices, statistical outliers against vendor history, and policy breaches like above-threshold amounts
Explanation Agent	Sujay Kumar	Turns each exception into a plain-language explanation a finance clerk can act on, citing the actual expected vs. found amounts
HITL Router	Riya Patil	Decides auto-clear vs. human queue using confidence, severity and amount at risk; builds the approval queue and records decisions back into state
Correspondence Agent	Toorubatla Upendra	Drafts a professional vendor dispute email citing the PO number and specific discrepancy — always a draft for review, never auto-sent